

bgphhypbn

April 12, 2024

0.0.1 Entrenamiento de DCGAN para proyecto 2.

Se utilizará la GAN llamada DCGAN, es una red neuronal convolucional. La arquitectura es la siguiente:

Generador:

Capa de entrada: Es una capa de entrada que toma una distribución de ruido aleatorio como entrada. Esta distribución de ruido suele seguir una distribución gaussiana o uniforme.

Capas ocultas convolucionales: Estas capas convolucionales toman la distribución de ruido y la transforman gradualmente en una imagen generada. Cada capa convolucional puede aumentar la dimensionalidad espacial de los datos mientras reduce la cantidad de canales.

Capa de salida: La capa de salida produce la imagen generada final. Generalmente, utiliza una función de activación como ReLU o una tangente hiperbólica para generar los píxeles de la imagen.

Discriminador:

Capa de entrada: Similar al generador, esta capa toma una imagen como entrada.

Capas convolucionales: El discriminador utiliza varias capas convolucionales para extraer características de la imagen de entrada.

Capa de salida: La capa de salida toma la representación de características extraída por las capas convolucionales y devuelve la probabilidad de que la imagen de entrada sea real o falsa.

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
[ ]: !pip install --upgrade tensorflow keras
```

```
[ ]: import keras
import tensorflow as tf

from keras import layers
from keras import ops
import matplotlib.pyplot as plt
import os
import gdown
```

```
from zipfile import ZipFile
```

El dataset utilizado fue obtenido en la página de [Kaggle](#). Se trata de imágenes tanto de jugadores de fútbol como personas que conforman el staff. Las imágenes son de tamaño 128x128 y están en formato JPG. Este dataset cuenta con 8599 imágenes que serán utilizadas para entrenar la DCGAN.

```
[ ]: !unzip '/content/drive/MyDrive/football.zip' -d /content/
```

```
[ ]: dataset = keras.utils.image_dataset_from_directory(
    "/content/football", label_mode=None, image_size=(64, 64), batch_size=32
)
dataset = dataset.map(lambda x: x / 255.0)
```

Found 8599 files.

```
[ ]: for x in dataset:
    plt.axis("off")
    plt.imshow((x.numpy() * 255).astype("int32")[1])
    break
```



Discriminador

```
[ ]: discriminator = keras.Sequential(
    [
```

```

keras.Input(shape=(64, 64, 3)),
layers.Conv2D(64, kernel_size=4, strides=2, padding="same"),
layers.LeakyReLU(negative_slope=0.2),
layers.Conv2D(128, kernel_size=4, strides=2, padding="same"),
layers.LeakyReLU(negative_slope=0.2),
layers.Conv2D(128, kernel_size=4, strides=2, padding="same"),
layers.LeakyReLU(negative_slope=0.2),
layers.Flatten(),
layers.Dropout(0.2),
layers.Dense(1, activation="sigmoid"),
],
name="discriminator",
)
discriminator.summary()

```

Model: "discriminator"

Layer (type) ↳Param #	Output Shape	
conv2d (Conv2D) ↳3,136	(None, 32, 32, 64)	↳
leaky_re_lu (LeakyReLU) ↳ 0	(None, 32, 32, 64)	↳
conv2d_1 (Conv2D) ↳131,200	(None, 16, 16, 128)	↳
leaky_re_lu_1 (LeakyReLU) ↳ 0	(None, 16, 16, 128)	↳
conv2d_2 (Conv2D) ↳262,272	(None, 8, 8, 128)	↳
leaky_re_lu_2 (LeakyReLU) ↳ 0	(None, 8, 8, 128)	↳
flatten (Flatten) ↳ 0	(None, 8192)	↳
dropout (Dropout) ↳ 0	(None, 8192)	↳

dense (Dense) (None, 1)
↳8,193

Total params: 404,801 (1.54 MB)

Trainable params: 404,801 (1.54 MB)

Non-trainable params: 0 (0.00 B)

Generator

```
[ ]: latent_dim = 128

generator = keras.Sequential(
    [
        keras.Input(shape=(latent_dim,)),
        layers.Dense(8 * 8 * 128),
        layers.Reshape((8, 8, 128)),
        layers.Conv2DTranspose(128, kernel_size=4, strides=2, padding="same"),
        layers.LeakyReLU(negative_slope=0.2),
        layers.Conv2DTranspose(256, kernel_size=4, strides=2, padding="same"),
        layers.LeakyReLU(negative_slope=0.2),
        layers.Conv2DTranspose(512, kernel_size=4, strides=2, padding="same"),
        layers.LeakyReLU(negative_slope=0.2),
        layers.Conv2D(3, kernel_size=5, padding="same", activation="sigmoid"),
    ],
    name="generator",
)
generator.summary()
```

Model: "generator"

Layer (type)	Output Shape
↳Param #	
dense_1 (Dense)	(None, 8192)
↳1,056,768	
reshape (Reshape)	(None, 8, 8, 128)
↳ 0	
conv2d_transpose (Conv2DTranspose)	(None, 16, 16, 128)
↳262,272	

leaky_re_lu_3 (LeakyReLU)	(None, 16, 16, 128)	└
↪ 0		
conv2d_transpose_1 (Conv2DTranspose)	(None, 32, 32, 256)	└
↪ 524,544		
leaky_re_lu_4 (LeakyReLU)	(None, 32, 32, 256)	└
↪ 0		
conv2d_transpose_2 (Conv2DTranspose)	(None, 64, 64, 512)	└
↪ 2,097,664		
leaky_re_lu_5 (LeakyReLU)	(None, 64, 64, 512)	└
↪ 0		
conv2d_3 (Conv2D)	(None, 64, 64, 3)	└
↪ 38,403		

Total params: 3,979,651 (15.18 MB)

Trainable params: 3,979,651 (15.18 MB)

Non-trainable params: 0 (0.00 B)

Definición de la GAN, juntando el discriminador y el generador.

```
[ ]: class GAN(keras.Model):
    def __init__(self, discriminator, generator, latent_dim):
        super().__init__()
        self.discriminator = discriminator
        self.generator = generator
        self.latent_dim = latent_dim
        self.seed_generator = keras.random.SeedGenerator(1337)

    def compile(self, d_optimizer, g_optimizer, loss_fn):
        super().compile()
        self.d_optimizer = d_optimizer
        self.g_optimizer = g_optimizer
        self.loss_fn = loss_fn
        self.d_loss_metric = keras.metrics.Mean(name="d_loss")
        self.g_loss_metric = keras.metrics.Mean(name="g_loss")

    @property
```

```

def metrics(self):
    return [self.d_loss_metric, self.g_loss_metric]

def train_step(self, real_images):
    # Sample random points in the latent space
    batch_size = ops.shape(real_images)[0]
    random_latent_vectors = keras.random.normal(
        shape=(batch_size, self.latent_dim), seed=self.seed_generator
    )

    # Decode them to fake images
    generated_images = self.generator(random_latent_vectors)

    # Combine them with real images
    combined_images = ops.concatenate([generated_images, real_images],
    ↪axis=0)

    # Assemble labels discriminating real from fake images
    labels = ops.concatenate(
        [ops.ones((batch_size, 1)), ops.zeros((batch_size, 1))], axis=0
    )

    # Add random noise to the labels - important trick!
    labels += 0.05 * tf.random.uniform(tf.shape(labels))

    # Train the discriminator
    with tf.GradientTape() as tape:
        predictions = self.discriminator(combined_images)
        d_loss = self.loss_fn(labels, predictions)
        grads = tape.gradient(d_loss, self.discriminator.trainable_weights)
        self.d_optimizer.apply_gradients(
            zip(grads, self.discriminator.trainable_weights)
        )

    # Sample random points in the latent space
    random_latent_vectors = keras.random.normal(
        shape=(batch_size, self.latent_dim), seed=self.seed_generator
    )

    # Assemble labels that say "all real images"
    misleading_labels = ops.zeros((batch_size, 1))

    # Train the generator (note that we should *not* update the weights
    # of the discriminator)!
    with tf.GradientTape() as tape:
        predictions = self.discriminator(self.
    ↪generator(random_latent_vectors))
        g_loss = self.loss_fn(misleading_labels, predictions)

```

```

        grads = tape.gradient(g_loss, self.generator.trainable_weights)
        self.g_optimizer.apply_gradients(zip(grads, self.generator.
↪ trainable_weights))

        # Update metrics
        self.d_loss_metric.update_state(d_loss)
        self.g_loss_metric.update_state(g_loss)
        return {
            "d_loss": self.d_loss_metric.result(),
            "g_loss": self.g_loss_metric.result(),
        }

```

```

[ ]: class GANMonitor(keras.callbacks.Callback):
    def __init__(self, num_img=3, latent_dim=128):
        self.num_img = num_img
        self.latent_dim = latent_dim
        self.seed_generator = keras.random.SeedGenerator(42)

    def on_epoch_end(self, epoch, logs=None):
        random_latent_vectors = keras.random.normal(
            shape=(self.num_img, self.latent_dim), seed=self.seed_generator
        )
        generated_images = self.model.generator(random_latent_vectors)
        generated_images *= 255
        generated_images.numpy()
        for i in range(self.num_img):
            img = keras.utils.array_to_img(generated_images[i])
            img.save("generated_img_%03d_%d.png" % (epoch, i))

```

Entrenamiento de la GAN con el dataset personalizado

```

[ ]: epochs = 50

gan = GAN(discriminator=discriminator, generator=generator,
↪ latent_dim=latent_dim)
gan.compile(
    d_optimizer=keras.optimizers.Adam(learning_rate=0.0001),
    g_optimizer=keras.optimizers.Adam(learning_rate=0.0001),
    loss_fn=keras.losses.BinaryCrossentropy(),
)

gan.fit(
    dataset, epochs=epochs, callbacks=[GANMonitor(num_img=10,
↪ latent_dim=latent_dim)]
)

```

Epoch 1/50

```

269/269          111s 323ms/step -
d_loss: 0.2325 - g_loss: 6.4681
Epoch 2/50
269/269          71s 263ms/step -
d_loss: 0.2073 - g_loss: 8.2717
Epoch 3/50
269/269          83s 268ms/step -
d_loss: 0.0953 - g_loss: 13.4840
Epoch 4/50
269/269          83s 271ms/step -
d_loss: 0.5094 - g_loss: 2.4636
Epoch 5/50
269/269          73s 271ms/step -
d_loss: 0.3381 - g_loss: 2.1143
Epoch 6/50
269/269          82s 272ms/step -
d_loss: 0.5671 - g_loss: 1.5214
Epoch 7/50
269/269          82s 273ms/step -
d_loss: 0.6264 - g_loss: 1.2138
Epoch 8/50
269/269          73s 271ms/step -
d_loss: 0.4386 - g_loss: 1.3631
Epoch 9/50
269/269          73s 271ms/step -
d_loss: 0.6789 - g_loss: 1.2438
Epoch 10/50
269/269          73s 272ms/step -
d_loss: 0.6162 - g_loss: 1.0661
Epoch 11/50
269/269          73s 271ms/step -
d_loss: 0.5815 - g_loss: 1.0307
Epoch 12/50
269/269          73s 271ms/step -
d_loss: 0.6538 - g_loss: 1.0502
Epoch 13/50
269/269          73s 271ms/step -
d_loss: 0.5390 - g_loss: 1.2672
Epoch 14/50
269/269          82s 272ms/step -
d_loss: 0.6235 - g_loss: 1.0322
Epoch 15/50
269/269          82s 273ms/step -
d_loss: 0.6543 - g_loss: 1.1024
Epoch 16/50
269/269          73s 271ms/step -
d_loss: 0.6554 - g_loss: 0.9187
Epoch 17/50

```



```

269/269          73s 271ms/step -
d_loss: 0.6685 - g_loss: 0.8846
Epoch 18/50
269/269          82s 272ms/step -
d_loss: 0.6831 - g_loss: 0.9171
Epoch 19/50
269/269          82s 273ms/step -
d_loss: 0.6103 - g_loss: 1.0200
Epoch 20/50
269/269          73s 271ms/step -
d_loss: 0.6279 - g_loss: 1.0129
Epoch 21/50
269/269          73s 271ms/step -
d_loss: 0.5825 - g_loss: 1.0746
Epoch 22/50
269/269          73s 271ms/step -
d_loss: 0.6002 - g_loss: 1.0556
Epoch 23/50
269/269          73s 271ms/step -
d_loss: 0.6083 - g_loss: 1.0404
Epoch 24/50
269/269          73s 271ms/step -
d_loss: 0.6654 - g_loss: 0.9314
Epoch 25/50
269/269          73s 271ms/step -
d_loss: 0.6194 - g_loss: 1.0045
Epoch 26/50
269/269          82s 272ms/step -
d_loss: 0.6582 - g_loss: 0.9918
Epoch 27/50
269/269          73s 271ms/step -
d_loss: 0.5421 - g_loss: 1.2651
Epoch 28/50
269/269          82s 272ms/step -
d_loss: 0.6262 - g_loss: 1.0056
Epoch 29/50
269/269          82s 273ms/step -
d_loss: 0.6638 - g_loss: 1.0034
Epoch 30/50
269/269          82s 273ms/step -
d_loss: 0.6198 - g_loss: 1.1296
Epoch 31/50
269/269          73s 272ms/step -
d_loss: 0.5627 - g_loss: 1.1923
Epoch 32/50
269/269          73s 271ms/step -
d_loss: 0.6081 - g_loss: 1.0598
Epoch 33/50

```

```

269/269          73s 272ms/step -
d_loss: 0.5986 - g_loss: 1.1126
Epoch 34/50
269/269          73s 272ms/step -
d_loss: 0.6205 - g_loss: 1.0024
Epoch 35/50
269/269          82s 272ms/step -
d_loss: 0.5981 - g_loss: 1.1747
Epoch 36/50
269/269          73s 271ms/step -
d_loss: 0.5806 - g_loss: 1.2837
Epoch 37/50
269/269          73s 271ms/step -
d_loss: 0.6035 - g_loss: 1.2493
Epoch 38/50
269/269          73s 271ms/step -
d_loss: 0.5524 - g_loss: 1.2170
Epoch 39/50
269/269          82s 273ms/step -
d_loss: 0.6210 - g_loss: 1.0341
Epoch 40/50
269/269          82s 273ms/step -
d_loss: 0.5825 - g_loss: 1.2155
Epoch 41/50
269/269          82s 273ms/step -
d_loss: 0.6495 - g_loss: 1.0771
Epoch 42/50
269/269          73s 271ms/step -
d_loss: 0.6238 - g_loss: 1.1163
Epoch 43/50
269/269          73s 271ms/step -
d_loss: 0.6203 - g_loss: 1.0844
Epoch 44/50
269/269          73s 272ms/step -
d_loss: 0.5488 - g_loss: 1.1360
Epoch 45/50
269/269          82s 273ms/step -
d_loss: 0.6079 - g_loss: 1.1080
Epoch 46/50
269/269          82s 273ms/step -
d_loss: 0.6497 - g_loss: 1.0560
Epoch 47/50
269/269          73s 271ms/step -
d_loss: 0.5977 - g_loss: 1.1085
Epoch 48/50
269/269          73s 271ms/step -
d_loss: 0.6189 - g_loss: 1.1370
Epoch 49/50

```

```
269/269          73s 272ms/step -
d_loss: 0.6251 - g_loss: 1.0323
Epoch 50/50
269/269          82s 272ms/step -
d_loss: 0.6159 - g_loss: 1.0649
```

```
[ ]: <keras.src.callbacks.history.History at 0x7b7ab054b1c0>
```

0.0.2 Veamos las imágenes generadas por la gan con 50 épocas.

```
[ ]: import matplotlib.pyplot as plt

num_images_to_generate = 10
latent_dim = 128

random_latent_vectors = tf.random.normal(shape=(num_images_to_generate,
↪latent_dim))

generated_images = generator.predict(random_latent_vectors)

# Muestra las imágenes generadas
plt.figure(figsize=(50, 50))
for i in range(num_images_to_generate):
    plt.subplot(1, num_images_to_generate, i + 1)
    plt.imshow(generated_images[i] * 0.5 + 0.5) # Escala las imágenes al rango
↪[0, 1]
    plt.axis("off")
plt.show()
```

1/1 2s 2s/step



0.0.3 Conclusiones

Como se puede observar, los resultados de la GAN no son perfectos, esto puede ser debido a que el número de épocas fue bajo. Al aumentar el número de épocas se pueden mejorar los resultados del generador. Sin embargo, la pérdida del generador fue alta, esto también podría mejorar con un mayor número de épocas. La pérdida del discriminador fue mejor, con una del casi 0.5, lo cuál indica que podría mejorar igualmente con más épocas.